

Disagreement

Edited by

Richard Feldman

and

Ted A. Warfield

OXFORD
UNIVERSITY PRESS

OXFORD
UNIVERSITY PRESS

Great Clarendon Street, Oxford OX2 6DP
Oxford University Press is a department of the University of Oxford.
It furthers the University's objective of excellence in research, scholarship,
and education by publishing worldwide in

Oxford New York

Auckland Cape Town Dar es Salaam Hong Kong Karachi
Kuala Lumpur Madrid Melbourne Mexico City Nairobi
New Delhi Shanghai Taipei Toronto

With offices in

Argentina Austria Brazil Chile Czech Republic France Greece
Guatemala Hungary Italy Japan South Korea Poland Portugal
Singapore Switzerland Thailand Turkey Ukraine Vietnam

Oxford is a registered trade mark of Oxford University Press
in the UK and in certain other countries

Published in the United States
by Oxford University Press Inc., New York

© the several contributors 2010

The moral rights of the author have been asserted

Database right Oxford University Press (maker)

Reprinted 2010

All rights reserved. No part of this publication may be reproduced,
stored in a retrieval system, or transmitted, in any form or by any means,
without the prior permission in writing of Oxford University Press,
or as expressly permitted by law, or under terms agreed with the appropriate
reprographics rights organization. Enquiries concerning reproduction
outside the scope of the above should be sent to the Rights Department,
Oxford University Press, at the address above

You must not circulate this book in any other binding or cover
And you must impose this same condition on any acquirer

ISBN 978-0-19-922608-5

Printed in the United Kingdom by
the MPG Books Group Ltd

For Ernie Sosa, despite occasional disagreements

Contents

<i>Contributors</i>	ix
Introduction	I
1. We're Right. They're Wrong <i>Peter van Inwagen</i>	10
2. Belief in the Face of Controversy <i>Hilary Kornblith</i>	29
3. Persistent Disagreement <i>Catherine Z. Elgin</i>	53
4. Rational Disagreement Defended <i>Earl Conee</i>	69
5. You Can't Trust a Philosopher <i>Richard Fumerton</i>	91
6. Peer Disagreement and Higher-Order Evidence <i>Thomas Kelly</i>	111
7. How to Disagree about How to Disagree <i>Adam Elga</i>	175
8. Epistemic Relativism and Reasonable Disagreement <i>Alvin I. Goldman</i>	187
9. The Moral Evil Demons <i>Ralph Wedgwood</i>	216
10. Disputing about Taste <i>Andy Egan</i>	247
<i>Index</i>	287

we have good reason to believe that individual experts are each highly reliable overall, and they should surely regard one another in this way. Their considered opinions should thereby be tempered when there is disagreement in the field. In philosophy, however, and the other areas of controversy we have been considering here, there is no such history of long-standing progress, and for that very reason we should not consider the experts in the field—including, of course, ourselves—to be highly reliable. The history of the field simply does not give any ground for credence in such a view. But it is then, for that very reason, that we must, in the end, withhold opinion on the issues under consideration.

I do not mean to suggest that we should stop thinking about these issues, or that thinking about them, and trying to work out tenable views, is not intellectually respectable. This is not my view at all. And, given the nature of human belief, I very much doubt that philosophers will stop forming views about the subjects they think about for so long and with such care. When we stand back, however, and reflect on our practice and on the beliefs that that practice generates, it seems to me that the history of our field makes epistemic modesty the only rational position available. We may hope that in trying to work out the views we are most sympathetic to that, in the long run, we may somehow contribute to an approach to philosophy that will look more progressive than any we have thus far seen. But, at the present time, we should all recognize that this is merely a hope, and that rational belief must be tempered by the facts about our current situation.

3

Persistent Disagreement

CATHERINE Z. ELGIN

1. The Problem

Some disagreements are epistemologically benign. One party is wrong and it is easy to determine which one. Sue, dimly recalling Longfellow's poem, believes that the American Revolution began on April 18, 1775; Sam, relying on his study of history rather than poetry, believes that it began on April 19, 1775. Dan believes that his flute is in tune; Dora, who has perfect pitch, believes that it is not. Maeve believes that, if smoking causes cancer, then she will not get cancer if she does not smoke; Mark, realizing that 'if p then q ' does not entail 'if not- p then not- q ', disagrees. Such differences of opinion are unthreatening. Sam, Dora, and Mark have excellent grounds for their beliefs, and excellent grounds for thinking that their opponents are mistaken. They should retain their beliefs, and be unmoved by their opponent's opinions. Disagreement *per se* does not jeopardize epistemic standing.

More problematic are cases in which opponents are, and consider themselves to be, epistemic peers. Then they have the same evidence, reasoning abilities, training, and background assumptions. If Fred and George are, and realize that they are, equally good at spelling, then Fred's firm belief that the proper spelling is 'ignomineous' when George believes it is spelled 'ignominious' should give them both pause. They should probably suspend judgment and consult a dictionary.

Although inconvenient, such short-term suspensions of judgment are relatively easily handled. The serious difficulty comes with persistent disagreement, where no easy or obvious resolution is available. The

evidence is equivocal. The evidence class contains misleaders, but there is no consensus about which the misleaders are. Should opponents suspend judgment in these cases too? Suppose two paleontologists, Jack and Jill, are epistemic peers who disagree about the fate of the Neanderthals. Jack believes that Neanderthals were an evolutionary dead end. Unable to compete, they simply died out. Jill believes that Neanderthals evolved into later hominids whose descendants are alive today. Because the issue is complex and the evidence is equivocal, they come to different conclusions about it. What should they (and we) make of their disagreement? In particular, should the fact that an epistemic peer disagrees with Jack have any effect on the epistemological status of his belief? Should Jack's knowledge of that fact have any effect?

Opinions diverge. Some philosophers, such as Richard Feldman, Hilary Kornblith, and David Christensen, contend that the existence of peer disagreement undermines one's grounds for belief.¹ If someone with the same evidence, training, background knowledge and reasoning abilities came to the opposite conclusion from Jack's, that is evidence that Jack's grounds are inadequate. Such philosophers think that epistemic agents should moderate their views in light of the disagreement. Others, such as Thomas Kelly and Richard Foley, maintain that it is reasonable for a thinker to retain his opinion in the face of disagreement with an epistemic peer.² They think that epistemic agents should be resolute. Both positions have unwelcome consequences.

2. Resoluteness

Advocates of resoluteness maintain that Jack should hold fast to his belief. To do otherwise would be spineless. Since Jack believes that his reasons are sufficient to support his conclusion, he thinks that Jill is wrong about

¹ Richard Feldman, "Epistemological Puzzles about Disagreement," in Stephen Hetherington (ed.), *Epistemology Futures* (Oxford: Oxford University Press, 2006), 216–36; Hilary Kornblith, "Belief in the Face of Controversy," Ch. 2, this volume; David Christensen, "Epistemology of Disagreement: The Good News," *Philosophical Review*, 116/2 (2007), 187–217.

² Thomas Kelly, 'The Epistemic Significance of Disagreement,' in Tamar Szabo Gendler and John Hawthorne, eds., *Oxford Studies in Epistemology*, i (Oxford: Oxford University Press, 2005), 167–96; Richard Foley, *Intellectual Trust in Oneself and Others* (Cambridge: Cambridge University Press, 2001), 83–131.

the Neanderthals. This is compatible with her being, and his recognizing her as, his epistemic peer. Everyone makes mistakes. So, although she is generally as good a paleontologist as he is, their disagreement is reason enough for him to conclude that in this case she is in error. Even though he cannot point to any flaw in her reasoning, Kelly maintains, Jack should take the disagreement to show that there must be a flaw. The mere fact that they disagree convicts Jill in Jack's eyes. Since they have the same evidence, her error must lie in how she reasons about the evidence. From Jack's point of view, the disagreement demonstrates that Jill is, in this case, not just wrong but irrational.

The situation is symmetrical. Jill, on Kelly's view, should be equally resolute. From Jill's perspective, Jack is being irrational. There currently is no tie-breaker. Such symmetry is disconcerting. It is evidently impossible to determine which of them is rational without determining which of them is right. So the fact that a belief is rational cannot function as a reason to think it is right. The stance of each party to such a disagreement is precarious. It is hard to be confident that one's belief is well founded while conceding that one's epistemic peer reasonably considers it irrational. Moreover, resoluteness seems to deprive epistemic agents of resources for correcting their mistakes. If Jack can so easily dismiss Jill's opinion, her disagreement gives him no reason to re-examine his own position, to seek further evidence, or to develop better methods of assessing the evidence.

Jack and Jill are experts in paleontology. So their disagreement occurs in a context in which they have what is, and what they recognize as, good evidence, and in which they have and recognize that they have, the capacity to reason responsibly about that evidence. But peer disagreements can occur at any level of expertise. Bill and Beth are epistemic peers who are woefully naive about economics. Bill believes that Liberia's dependence on US currency weakens the dollar. Beth believes that it does not. They appeal to the same sparse and dubiously relevant considerations to justify their conclusions. Once their disagreement focuses attention on the inadequacy of the reasons they can offer to support their positions, it becomes clear that neither has a clue whether Liberia has any effect on the value of the US dollar. Plainly, they should suspend belief. Evidently, a threshold of competence has to be reached before resoluteness is even remotely reasonable. Only if epistemic peers are cognitively competent with regard

to the topic under dispute is it plausible that they should retain their belief in the face of disagreement.

3. Moderation

Feldman maintains that epistemic peers who disagree should suspend belief. Christensen maintains that each of them should moderate his degree of belief, although not perhaps to the point of suspending belief entirely. For reasons of simplicity, I will discuss only Feldman's position, but my points extend in obvious ways to Christensen's.

In cognitive contexts, Feldman notes, it is always open to an epistemic agent to suspend belief.³ She should do so whenever she recognizes that her grounds are inadequate. Jill's epistemic peer disagrees with her about the fate of the Neanderthals. She can find no fault with his reasoning. This, Feldman maintains, provides Jill with evidence that her grounds are inadequate. So she should suspend belief. The epistemic situation is symmetrical. Jack should suspend belief too. Symmetry, an unattractive feature of resolute theories, is unproblematic for advocates of moderation, for it leads to a convergence of opinion.

Although suspending belief in such cases may seem reasonable, it pushes in the direction of skepticism. Wherever issues are complicated, if there are epistemic peers, they are apt to disagree. Thus there are apt to be vast areas of inquiry where belief is to be suspended. Moreover, whether one happens to have an epistemic peer seems to be utterly contingent. This suggests that we should consider possible peer disagreement as well. Should Harry suspend belief because he recognizes that if he had any epistemic peers, some of them would disagree?

One might think not. We may imagine that thinkers with no epistemic peers are geniuses—people like Einstein or Darwin. The fact that they have no peers is reason to believe that they at least need not suspend belief. Their reasons and reasoning powers are so strong that no one with those reasons and reasoning powers could disagree. But in the recent debates about disagreement 'epistemic peer' is defined quite narrowly. It requires having *the same* evidence and reasoning abilities. So it is not surprising if an

ordinary person lacks epistemic peers with respect to a particular, mundane issue. If Jen and Jon have even slightly different relevant reasoning abilities or evidence pertaining to the causes of the Civil War, they are not epistemic peers with respect to the subject. Given the vicissitudes of education and abilities, and the idiosyncrasies of evidence-gathering, ordinary epistemic agents are apt to have few epistemic peers. But, if the only reason that Harry does not face peer disagreement about a particular issue is that he happens to have no epistemic peers with respect to that issue, then the absence of disagreement is fortuitous. It hardly puts his belief about the fate of the Neanderthals on a stronger epistemic footing than Jack's. In these discussions 'epistemic peer' is an idealization. To decide when and how real disagreement should affect real epistemic agents we need either to construe 'epistemic peer' more generously or to take seriously possible as well as actual epistemic peers. Rather than introduce possible peers, I will construe peer disagreement more generously, so that epistemic peers are those who have pretty much the same relevant evidence, reasoning powers, training, and background information. This choice is a matter of expository convenience. My discussion could equally be cast in terms of actual and possible peers, where the standard conception of an epistemic peer is used.

Either way, the unwelcome consequence of Feldman's view is that it recommends suspending judgment in a vast number of cases. We would be forced to skepticism about such things as the fate of the Neanderthals, the causes of the Civil War, the problem of free will, the Kennedy assassination, and so on. On Christensen's view we would not necessarily be forced to skepticism, but disagreements about such matters would require us to moderate the strength of our beliefs. The fact that some of Joe's epistemic peers are incompatibilists means that he should not strongly believe that free will and determinism are compatible.

4. Hyperresoluteness

So far, we have considered the cases where peers disagree, and the choices we have entertained are between moderating one's views and standing firm—roughly, between being spineless and being stubborn. Epistemic agents who moderate their beliefs in the face of disagreement seem spineless,

³ Feldman, "Epistemological Puzzles about Disagreement," 229.

abandoning their convictions as soon as a serious challenge appears on the scene. Resolute epistemic agents seem stubborn, simply insisting that there must be something wrong with their opponent's reasoning, since there is plainly nothing wrong with their own. To see the way out of this predicament, let us look at a more extreme case raised by Peter van Inwagen.⁴ Call this the David Lewis Problem.

David Lewis believed that infinitely many possible worlds exist, each of them just as real as the actual world.⁵ There is no denying that he believed this. Moreover, there is no denying that he was incredibly smart, philosophically gifted, and intellectually responsible. He examined the arguments for and against his position with enormous care. It is no false modesty for me to say that David Lewis was a far better philosopher than I am. Nevertheless, I think he was wrong. I cannot refute his position; it is admirably well defended. But, despite Lewis's intelligence and arguments, I do not believe that there exist real possible worlds, consisting of material objects and inaccessible from the actual world.

I believe that the only world is the actual world. I think that my belief is reasonable. But David Lewis thought otherwise. He was not my epistemological peer; he was my epistemological superior. So should I not revise my opinion to agree with him? After all, if a knowledgeable physicist tells me that, despite what I think, electrons are not material particles, but clouds of energy, I revise my belief to accord with hers. So in some cases, at least, it seems epistemically reasonable to defer to my epistemic superiors. Is my disagreement with Lewis different?

Van Inwagen's answer is similar to Kelly's. In explaining why he thinks it is reasonable to retain philosophical convictions with which Lewis disagrees, he says:

I suppose my best guess is that I enjoy some sort of philosophical insight [with respect to these issues] that, for all his merits, is somehow denied to Lewis. And this would have to be an insight that is incommunicable—at least *I* don't know how to communicate it—for I have done all that I can to communicate it to Lewis,

⁴ Peter van Inwagen, "It is Wrong Everywhere, Always, and for Anyone to Believe Anything upon Insufficient Evidence," in Eleanor Stump and Michael J. Murray (eds.), *Philosophy of Religion: The Big Questions* (Oxford: Blackwell, 1999), 274. I consider only one of the disagreements with Lewis that van Inwagen mentions.

⁵ David Lewis, *Counterfactuals* (Cambridge, MA: Harvard University Press, 1973); *On the Plurality of Worlds* (Oxford: Blackwell, 1986).

and he has understood perfectly everything I have said, and he has not come to share my conclusions.⁶

Van Inwagen thinks that the disagreement shows that Lewis must be mistaken, even though he cannot say what the mistake is. Kelly would add that Lewis, being wrong, must be irrational.

I cannot speak for van Inwagen. But, speaking for myself, I think it is exceedingly unlikely that I enjoy any sort of philosophical insight that Lewis lacked (except, perhaps for the utterly question-begging insight that I am right and he is wrong, which even if true is utterly question-begging). Nor can I conclude, as advocates of resoluteness think I should, that in this case Lewis's reasoning is flawed. The position is amply, publicly, brilliantly defended. The number of able philosophers who cannot find a defect in the argument is legion. And, in response to an endless barrage of criticisms and incredulous stares, Lewis re-examined his position often. Granted, there may nevertheless be an extremely subtle flaw in Lewis's reasoning. But, on the available evidence, it is sheer hubris to insist that there must be. I do not believe that Lewis was being irrational. Should I, conceding Lewis's epistemic superiority in metaphysics, endorse realism about possible worlds? If not, should I at least follow Feldman's advice and suspend belief? Although I am not Lewis's *peer*, I might be close enough to a peer for it to be epistemically permissible for me to suspend belief rather than going over to his side. But even this seems excessively open-minded. I do not even think that Lewis might be right on this matter. Am I being hyperresolute? Should I be?

5. The Solution

Luckily, there is an easy solution to the David Lewis Problem. Unfortunately, it simply unmasks a further problem. For, if we accept this solution, as I think we should, we must conclude that recent debates about the epistemic consequences of disagreement rest on a mistake.

The solution is this: despite the fact that Lewis's position is brilliantly constructed, admirably defended, and beautifully argued, I find it incredible.

⁶ Van Inwagen, "It is Wrong Everywhere, Always, and for Anyone," 274.

I simply cannot believe it. Since 'ought' implies 'can,' that I cannot believe it entails that it is not the case that I ought to believe it. And that I cannot believe that it might be true entails that it is not the case that I should suspend belief or lower my degree of belief that the only real world is the actual world. It is philosophically interesting and perhaps troubling that a position I find utterly incredible admits of such a strong defense. But my belief in a unique world is not in jeopardy.

One might think that this solution is available in extreme cases like the David Lewis Problem but not in ordinary cases like the ones we started with. That there are infinitely many real possible worlds is incredible; that the Neanderthals were an evolutionary dead end seems not to be. Jonathan Adler argues otherwise.⁷ Belief is responsive to evidence. Given a body of evidence, there is no choice about what to believe. So, even if it is not a priori incredible that the Neanderthals were an evolutionary dead end, when Jill surveys the evidence she finds it incredible, given that evidence, that the Neanderthals were an evolutionary dead end. In light of the evidence, she cannot believe it. Different epistemic agents might assess the evidence differently and so come to different beliefs. But this is not a matter of choice. They come to different beliefs because the evidence affects them differently.

Belief is not voluntary. Belief aims at truth in the sense that a belief is defective if its content is not true. If believing were something we could do or refrain from doing at will, the connection to truth would be severed. If Jack could believe that Neanderthals were an evolutionary dead end just because he wanted to, then his believing that Neanderthals were an evolutionary dead end would not amount to his thinking that 'Neanderthals were an evolutionary dead end' is true. For nothing about the fate of the Neanderthals is affected by what he wants. This is Bernard Williams's point.

If I could acquire a belief at will, I could acquire it whether it was true or not; moreover, I could know that I could acquire it whether it was true or not. If in full consciousness, I could will to acquire a 'belief' irrespective of its truth, it is unclear that . . . I could seriously think of it as a belief, i.e. as something purporting to represent reality.⁸

⁷ Jonathan Adler, *Belief's Own Ethics* (Cambridge, MA: MIT Press, 2002), 55–64.

⁸ Bernard Williams, "Deciding to Believe," in *Problems of the Self* (Cambridge: Cambridge University Press, 1973), 148.

Although evidence or other epistemic considerations can move me, since belief is not voluntary, my reaction is not anything I *do*. Being responsive or unresponsive to evidence, argument, or peer pressure is something that happens to me. This means that I cannot follow Kelly's recommendation that I hold fast to my belief in the face of peer disagreement or Feldman's recommendation that I suspend belief. How peer disagreement affects my belief is not up to me. I may find myself with a belief suspended as a result of evidence, argument, testimony, or disagreement; or I may find my belief unmoved by evidence, argument, testimony, or disagreement. But my response is not under my control. Debates about whether I should suspend belief in the face of peer disagreement are wrong-headed. They are like debates about whether I should be less than 6 feet tall. I do not have any choice.

6. Assessment

Although involuntariness does not exempt responses to disagreement from assessment, it affects the sort of assessment they are subject to. Plenty of things that are beyond our control are subject to assessment—the weather, for instance. It is a miserable day today; it would have been preferable if the rain had held off until after the parade. This is a perfectly respectable assessment that does not impute fault. If belief is not subject to direct voluntary control, then assessing someone's reaction to disagreement is similar to assessing the weather. 'It is (or is not) regrettable that Jack suspended belief' is like 'It is (or is not) regrettable that it rained on the parade.' The epistemological issue under dispute, then, is whether a better constellation of beliefs results when one suspends belief in the face of peer (or superior) disagreement or when one resolutely retains one's belief. The locus of assessment is the constellation of beliefs, not the actions of the believer.

If such assessments are assessments of doxastic rationality, then assessments of doxastic rationality are assessments of what happens to us cognitively. In that case, they are like 'Jim is smart' and 'Joan is creative.' If suspending belief (or retaining belief) is cognitively good in cases of disagreement with epistemic peers or epistemic superiors, and 'rationality' is the predicate that characterizes that sort of goodness, then 'He is rational' is the same sort

of praise as 'She is smart.' Both characterize cognitively valuable attributes that their subjects happen to possess. Neither characterizes an attribute that is under their subjects' control.

This construal of rationality would allow us to characterize those who respond correctly to peer disagreement as rational. But it would have the consequence that doxastic rationality and practical rationality diverge. Practical rationality presupposes control. Actions are voluntary, and the rationality of an action depends on what an agent voluntarily does, given her beliefs, desires, preferences, and so on. Behaviors such as sneezes, spasms, and snores are exempt from assessment as rational or irrational, because they are involuntary. If practical rationality depends on what we do and doxastic rationality depends on what happens to us, then the concept of rationality bifurcates. The term 'rational' indicates something quite different when applied to beliefs and when applied to action. This could be so, but it is an awkward result. It raises the question why we use the same term for two such different phenomena.

7. Indirect Control

Since beliefs are not voluntary, an epistemic agent cannot, even through judicious assessment, bring it about that she retains, or lowers her degree of belief, or suspends belief in the face of a disagreement. She may, however, be able to affect her responses indirectly. Pascal recognizes this in his discussion of the wager.⁹ He does not think that one could come to believe that God exists simply by appreciating that it would be prudent to believe that God exists. But he thinks that appreciating that it would be prudent to believe that God exists gives a person reason to put himself in a position to improve his prospects of acquiring the belief that God exists. By engaging in religious practices, interacting with religious people, and avoiding irreligious people and situations, Pascal maintains, a person maximizes his prospects of being moved by factors that foster the belief that God exists. Education has a similar effect. By learning about the cognitive force of evidence, argument, and expertise, students can be put in a position to be moved by considerations of one sort or another. And, as both Pascal

and the educators recognize, epistemic agents can learn to appreciate why it might be worthwhile to maximize their prospects of forming, retaining, revising, and rejecting beliefs of different kinds. Arguably then, we are rational vis-à-vis our belief that *p*, not directly because we are moved by the evidence for *p*, but because we properly put ourselves in a position to be so moved. If so, doxastic rationality is a sort of practical rationality. It applies to strategies for acquiring beliefs, not to beliefs themselves.

If this is so, then the real issue about the epistemological implications of disagreement is not whether an epistemic agent should retain or revise a belief in the face of peer disagreement. It is whether she should put herself in a position to be moved by such disagreement or put herself in a position to stand fast in the face of it. Either option would be, to a significant extent, a consequence of education. If an epistemic agent learns to appreciate the merit of her opponent's position or the value of his insights, she might find herself moderating her views when they conflict with his. If she learns the perils of skepticism and spinelessness, she might find her resistance to epistemic peer pressure strengthened. Such responses are effects of cognitive character formation. The question then is what sort of character we ought to form. There is, of course, no guarantee that our beliefs will respond as we hope they will. But by subjecting ourselves to the right influences, we maximize our prospects.

Such cognitive character formation is epistemically valuable. It is cognitively worthwhile to be able to appreciate how evidence, argument, and expertise bear on the tenability of a thesis, even where we cannot believe its conclusion. Although Hume recognizes that we cannot long sustain a skeptical attitude, he thinks that the arguments that lead to skepticism are important, for they reveal the flimsiness of our grounds for belief.¹⁰ Even if we cannot help but believe, we are better off knowing that our cognitive house is built on sand.

8. From Belief to Acceptance

If we retain the focus on belief, the relation to voluntariness and epistemic responsibility is distant at best. Let us consider an alternative. L. Jonathan

⁹ Blaise Pascal, "The Wager," in *Pensées*, ed. Roger Ariew (Indianapolis: Hackett, 2005), 212–14.

¹⁰ David Hume, *An Enquiry Concerning Human Understanding* (Indianapolis: Hackett, 1977), 25–37.

Cohen distinguishes between belief and acceptance.¹¹ To believe that p is to feel that p is so. To accept that p is to adopt a policy of being willing to treat p as a premise in inferences or as a basis for action. Let us modify this slightly and say that to accept that p is to adopt a policy of being willing to treat p as a premise in assertoric inference or as a basis for action where our interests are cognitive. The reason for the restriction to assertoric inference is to screen off premises used in reductios. The reason for the restriction to cognitive interests is that a premise accepted because it fosters non-cognitive ends—because it is consoling or amusing, for example—is epistemologically irrelevant.

Acceptance and belief are distinct. I can believe that p , and yet refuse to use it as a premise in inference or a basis for action. I might, for example, consider my evidence inadequate. Suppose I have what I consider an unfounded fear of frogs. (I was frightened by a frog at an early age and never quite got over it.) In that case, even though I cannot help but feel that frogs are dangerous, I refrain from using ‘frogs are dangerous’ in any cognitively serious inference. Nor do I act on my fear of frogs. I do not, for example, arm myself against them. Analogously, I can accept that p , even though I do not believe that p . I accept that frogs are not dangerous, when I include ‘frogs are not dangerous’ among the premises I am prepared to use when deliberating about the perils of summer camp, and when I allow my toddler to wade in a frog pond.

I suggest that the epistemic issue raised by peer disagreement is best seen as an issue pertaining to acceptance rather than belief. Accepting is something we do; it is an action. Hence it is voluntary. So we can, at will, continue to accept that p in the face of peer disagreement, or suspend acceptance that p . On this reading, if moderationists are right, peer disagreement should affect what we accept. It should influence the inferences we are prepared to make and actions we are prepared to perform. If advocates of resoluteness are right, it should not.

It might seem that this brings us back to where we started, or close enough to make no difference. The dispute, as I reconstrue it, concerns accepting for cognitive purposes rather than believing. But the issues are the same; the considerations favoring each side are the same; and we are no closer to finding a decisive reason to favor either holding fast to one’s

cognitive commitment or suspending that commitment in the face of peer disagreement.

I do not think this is so. The shift from belief to acceptance reconfigures the epistemic terrain. As Feldman points out, when our goals are purely cognitive, suspending judgment is always an option. He argues that it is the option we should exercise when faced with a disagreement with our epistemic peers. He recognizes that this pushes in the direction of skepticism, but considers such skepticism plausible. If we focus exclusively on the question whether to affirm p or to affirm that not- p or to affirm neither, suspending belief in the circumstances seems the safest thing to do. But the switch to acceptance highlights the fact that all three options have costs. To suspend acceptance vis-à-vis p is to adopt a policy of refraining from using either p or not- p as a premise in assertoric inferences or as a basis for action. This is a cognitively impoverishing stance. We have fewer premises available to reason with. If, for example, Jack suspends acceptance of ‘Neanderthals were an evolutionary dead end’, he cannot use it as a premise. If ‘Neanderthals were an evolutionary dead end’ were the only problematic claim in the neighborhood, suspending acceptance might be reasonable. But in cases like this, where evidence is sparse and equivocal, peers are apt to disagree about a host of related issues.¹² The extent of Neanderthal tool-making, the structures of their communities, their level of cognitive development, their resistance to disease are all controversial matters. What a paleontologist thinks about one of these issues is apt to be enmeshed with his position on others. There are apt to be complicated patterns of agreement and disagreement across the community. Jack may disagree with Jill about some matters and with Jen about others and with Joe about yet others. To suspend acceptance of all of them leaves the paleontological community with few premises about their subject matter, yielding a flimsy and moth-eaten fabric of cognitive commitments. It is not clear how they should reason about the paleolithic period, if they can deploy only premises about which no peer disagrees. Suspending acceptance in the face of peer disagreement in cases like this would probably be excessively cognitively costly.

¹¹ L. Jonathan Cohen, *An Essay on Belief and Acceptance* (Oxford: Oxford University Press, 1992).

¹² Adam Elga maintains that, if they disagree on a host of related issues, they do not share enough to be epistemic peers. This strikes me as implausible. See Adam Elga, “Reflection and Disagreement,” *Noûs*, 41/3 (2007), 478–502, and Hilary Kornblith, Ch. 2, this volume, for an argument to the contrary.

Whether, in the face of peer disagreement, to continue to accept that p , to come to accept that not- p , or to suspend acceptance vis-à-vis p is a practical question. It depends on what we gain and what we lose under each alternative, and what costs are worth paying. Which premises must we abandon? How central are they? Again, this might seem to collapse into the original debate, or something close to it. Standardly, epistemologists take it for granted that, if the evidence favors p strongly enough, it is rational to accept or believe that p ; if it favors not- p strongly enough, it is rational to accept or believe that not- p . If the evidence is about equally balanced, it is rational to suspend judgment. The standard of rationality is utterly general; everyone with the same evidence should respond the same way. The issue raised by peer disagreement is this: if Jack takes the evidence to support p strongly enough, and Jill, his epistemic peer, takes it to support not- p strongly enough, then Jack either needs somehow to discredit his epistemic peer, or take the disagreement to show that the evidence does not support either p or not- p strongly enough to justify belief or acceptance. Insisting that we count the costs does not seem to settle the issue.

9. The Value of Disagreement

The objectionable feature of resoluteness is the requirement that, to hold fast to their opinions, peers must construe each other as irrational. This is awkward, since they cannot point to any flaws in their opponents' reasoning and since they recognize that their opponents construe them as irrational on the same grounds. Moreover, if we think about actual cases of sustained peer disagreement, the charge of irrationality seems unfounded.

Jury deliberations are a familiar example. Some jurors think the defendant is guilty; others think she is not guilty. They all have the same evidence and, let us assume, the same reasoning powers. They disagree because they assess the evidence differently. It is clear to everyone that some of the evidence offered at the trial is misleading. A member of an opposing gang placed the defendant near the scene of the crime. A member of her own gang said that she was across town. The jurors disagree about which witness is reliable. Some doubt the first, since she bears the defendant a grudge; some doubt the second witness, since she seems like the sort who would lie to help her friend. Neither of the witnesses comes off as a stellar character. Jury

members might also disagree about the weight that attaches to various bits of evidence. How significant is it that the weapon was never found? How directly does the statistical evidence bear on a case like this? What should they make of the absence of fingerprints? It is not obvious that any of the jurors is irrational. If they cannot resolve their disagreement, the result will be a hung jury, the judicial equivalent of suspension of judgment. But, it seems, each juror could rationally retain his belief, while recognizing the rationality of his opponent's.

The drive to consensus in jury deliberations derives from their role in criminal trials. A hung jury is a disappointing outcome, for obvious practical reasons. But it is not clear that this point generalizes. Persistent disagreement in science or philosophy is not obviously a bad thing.

Consider the disagreement between materialists and dualists in the philosophy of mind. Materialists accept that whatever is is material; dualists accept that there are irreducibly mental entities or processes as well as irreducibly physical entities and processes. Each side can point to some conspicuous explanatory successes. But each side faces serious difficulties. Either there are outstanding problems that it cannot solve, or the solutions it offers seem inelegant, strained, and ad hoc. What is worse, the serious problems that each faces seem straightforwardly handled by the other. The dualist has a problem explaining the causal link between the mental and the material; the materialist can simply maintain that the connection is straightforward physical causality. The materialist has a problem accounting for qualia and what-is-it-like-ishness; the dualist takes these features to be distinctive marks of the mental. This is all familiar.

The standard epistemological view would maintain that in such a case everyone should go with the balance of evidence. There is a tipping point. Until the evidence reaches the tipping point, everyone should suspend acceptance. Once it is reached, everyone should accept whichever side the evidence favors. But, as Philip Kitcher argues, it is not obvious that our cognitive objectives are best achieved by everyone's marching in lock step to the same conclusion.¹³ When the reasons favoring each side of a dispute are sparse or exceedingly delicate, or the evidence is equivocal, or each side can solve important common problems that the other cannot, it may be better for the epistemic community as a whole that some of its members

¹³ Philip Kitcher, "The Division of Cognitive Labor," *Journal of Philosophy*, 87 (1990), 5–22.

continue to accept each position. In that case, materialists can in good conscience continue to accept materialism. Dualists can in good conscience continue to accept dualism. Agnostics can suspend judgment. Each group then can draw on a different range of commitments for premises in their reasoning and as a basis for their actions. By developing their positions, they put them to the test. Arguably, the only way we will ever find out whether materialism can solve the hard problem of consciousness is for materialists wholeheartedly to accept materialism and push it to its limits.

This position does not require denying that the overarching epistemic goal is accepting or believing only what is true. It simply notes that, where there is a significant chance that my opponent's view is true, if I want to believe only what is true, I would be well served by not foreclosing inquiry prematurely. If I can recognize that my opponent is rational and might (although I strongly doubt it) be right, then I have reason to hope that she retains her position, develops it, and either comes (as I believe she will) to see the error of her ways or (however unlikely) develops an argument that will demonstrate to me the error of mine. A convinced materialist then has sound epistemic reasons to tolerate dualism.

Such tolerance has limits, though. As we have seen, the mere fact of disagreement is not enough to make each party's position creditable. Nor is the mere fact of peer disagreement. Tolerance of disagreement is epistemically valuable only when the disagreement is among parties who have sufficient expertise in an area that their opinions are individually worth accepting, and where the evidence at hand is equivocal. When these conditions are met, a community of inquiry may best be served if epistemic peers resolutely reason and act on opinions about which other equally competent inquirers disagree. In such cases, peers who disagree have reason to consider each other wrong but not irrational. Perhaps in the fullness of time the disagreement will be resolved. That remains to be seen.

4

Rational Disagreement Defended

EARL CONEE

1. Introduction

We are rational only if we heed the dictates of reason. We heed our epistemic reasons by taking the doxastic attitudes that align with their support. If the bearing of our reasons on the truth-value of a proposition is unequivocal, then just one doxastic attitude is epistemically rational for us to take toward the proposition. Rational uniqueness principles imply that our epistemic reasons determine the doxastic attitudes that are uniquely epistemically rational for us to take.¹ The most defensible of these principles are restricted to cases of extreme duplication among those whom the principles imply to have the same justified attitudes. These epistemic peers must have identical total evidence, or they must have identical reasons, reasoning abilities, and inclinations to reason. Such restrictions are apparently excessive. It seems equally credible to require only a sharing of all that is rationally relevant to a proposition's truth-value.

Are some such relaxed uniqueness principles correct, or can thoroughly reasonable people share everything relevant to a disputed proposition and still rationally disagree? It will be argued here that the relaxed principles are not correct. Sharing the relevant reasons allows that the parties involved

¹ Our epistemic reasons could determine which are our epistemically rational attitudes toward a given proposition, while the reasons allow that more than one attitude is epistemically rational. Rational uniqueness principles go farther. They imply that a single doxastic attitude is rationally allowed by the reasons. (I am grateful to Tim Williamson for emphasizing this distinction in a discussion of an earlier draft of this chapter.) This difference does not affect the arguments of the chapter. They oppose a uniqueness principle by arguing that differing attitudes are uniquely rational.